

Software 3.0: Fine-tuning

Saeed Siddik

Assistant Professor

IIT University of Dhaka

Developing Alware systems

- Prompt Engineering
- **Fine-tuning**
- Retrieval Augmented Generation (RAG)

Fine tuning

- adapting a pre-trained AI-Model to perform better on a specific task or with a domain-specific dataset.
- The goal is to adapt the model's general knowledge to a specialized task, domain, or style.

Why Fine-Tune? The PM's View

Training from Scratch

- Requires **massive** datasets
- Extremely high compute cost
- Very high risk of failure
- Impractical for 99% of software projects

Fine-Tuning

- Leverages existing knowledge
- Requires a smaller, curated dataset
(100s-1000s of examples)
- Drastically lower cost and time
- Faster path to a specialized model

When NOT to fine-tune

- If non-compliant with privacy/regulatory constraints and model contains sensitive pretraining data you cannot modify.
- When task requires radically different architecture/inductive biases.
- When you must guarantee zero change from pretrained weights (use inference-time adapters instead).
- When compute budget or latency demands prohibit larger models like considering smaller models or distilled variants.

Types of Fine-Tuning (Part 1)



1. Full Fine-Tuning

Updates all weights of the pre-trained model. It's the most thorough method but is computationally expensive and creates a full-size copy of the model.



2. PEFT


Parameter-Efficient Fine-Tuning. Updates only a **small subset** of parameters (or adds new ones). It's much faster, cheaper, and creates tiny "adapter" files (e.g., 50MB vs 15GB).

Types of Fine-Tuning (Part 2): PEFT

Popular Parameter-Efficient Fine-Tuning Methods

- ✿ **LoRA (Low-Rank Adaptation):** Injects small, trainable "adapter" matrices into the model. The original model weights are frozen. This is the most popular and effective method.
- ‡ **QLoRA (Quantized LoRA):** An optimization of LoRA. Uses quantization to reduce the model's memory footprint, allowing fine-tuning of huge models on smaller, cheaper GPUs.

Project workflow of Fine-Tuning Process



Example: Task and data

- **Task:** sentiment classification for customer reviews (3 classes: positive, neutral, negative).
- **Pretrained model:** BERT-base (or a Transformer encoder).
- **Dataset:** 10k labeled reviews, 80/10/10 train/val/test.
- **Evaluation metric:** accuracy + F1 (macro).

Hands-On: What it Looks Like

Key Libraries (Python)

In a real project, you'd use open-source libraries to handle the complexity.

- **Hugging Face transformers:** To load the pre-trained model.
- **Hugging Face peft:** To easily apply PEFT methods like LoRA.
- **Hugging Face datasets:** To load and process your custom data.

Conceptual Code

```
# 1. Import libraries
from transformers import AutoModel, Trainer
from peft import get_peft_model, LoraConfig

# 2. Load base model
model = AutoModel.from_pretrained("meta-llama/Llama-3-8B")

# 3. Define PEFT (LoRA) config
lora_config = LoraConfig(r=8, lora_alpha=16)


# 4. Apply PEFT to the model
peft_model = get_peft_model(model, lora_config)

# 5. Run training
trainer = Trainer(model=peft_model, train_dataset=my_data)
trainer.train()
```

Low-Rank Adaptation

- LoRA (Low-Rank Adaptation) is a parameter-efficient fine-tuning method that freezes the original model weights and injects small trainable low-rank matrices into specific layers (usually attention layers).
- During training, only these small matrices are updated, not the full model.

LoRA: Low-Rank Adaptation



What is QLoRA?

- QLoRA extends LoRA by running the base model in 4-bit quantized form, drastically reducing memory, while still training LoRA adapters in higher precision (NF4/FP16).

LoRA vs QLoRA

Aspect	LoRA	QLoRA
Base model precision	FP16 / BF16	4-bit quantized
Memory savings	High	Very high ($3\times$ – $5\times$ more)
Accuracy	Near full	Nearto LoRA
Training speed	Good	Slightly slower (dequantization cost)
Hardware needs	Multiple GPUs for 13B+	Single 24GB GPU for 65B (with tricks)
Use case	Mid-size models	Very large LLMs

Building an "Alware" Project with LoRA



1. Data Pipeline

The project is the data. Your primary focus must be on robust data collection, cleaning, and versioning. Garbage In, Garbage Out.



2. Experiment Loop

The 'build' phase is an experiment. You must track the base model, dataset version, hyperparameters, and metrics for **every single run**.



3. Model Deployment

The output isn't just code; it's a model artifact (e.g., a 50MB LoRA adapter). This needs its own deployment and serving strategy.

Project management for fine-tuned AI systems

- **Phases:** discovery → data collection → modeling → integration → testing → deployment → monitoring → maintenance.
- **Roles:** Project Manager, Software Engineer, ML Engineer, Data Engineer, Annotator, QA, DevOps, UX.
- **Deliverables:** dataset spec, baseline, model artifacts, evaluation report, monitoring plan, runbooks.

Management Challenges (Part 1)

Data & Cost Management

Acquiring high-quality, **labeled** data is the #1 bottleneck and cost driver.

Compute (GPU) time is your most expensive resource and must be tracked like any other project cost.

The Versioning Nightmare

You must track **all four** components for reproducibility:

Base Model version (e.g., Llama-3-8B)

Dataset version (e.g., tickets_v1.1)

Training Code version (Git commit)

Resulting Model Artifact

Management Challenges (Part 2)

Evaluation is Hard

How do you *know* it's better?

Accuracy isn't enough. You must test for bias, fairness, robustness, and specific failure modes. Define success metrics before you start.

Model Drift & Forgetting

The real world changes. Your model's performance will 'drift' (decay) over time. It may also 'forget' general knowledge after being over-specialized (known as catastrophic forgetting).

Risk management Challenges

Data risk: low quality or insufficient labels.

Technical risk: model fails at scale, latency issues.

Ethical/legal: bias, privacy infractions.

Operational: costly retraining, lack of monitoring.

Mitigation: pilot studies, audits, staging, KPIs & SLAs.

End of Fine Tuning